

Lei Wang

+1 (424) 402-6563 | leiwang1@ufl.edu | [lei-wang-link.github.io](https://github.com/lei-wang-link) | Gainesville, FL

EDUCATION

University of Florida

Ph.D. Candidate in Electrical and Computer Engineering | GPA: 4.00/4.00

Gainesville, FL

Aug. 2022 – May. 2027 (Expected)

University of California, Los Angeles

M.S. in Electrical and Computer Engineering | GPA: 3.93/4.00

Los Angeles, CA

Sep. 2021 – Jun. 2022

University of Electronic Science and Technology of China

*B.E. in Electronic and Information Engineering | GPA: 3.97/4.00 | **National Scholarship***

Chengdu, China

Sep. 2016 – Jul. 2020

PROFESSIONAL EXPERIENCES

PayPal, Inc.

Machine Learning Scientist Intern

San Jose, CA

May. 2025 – Aug. 2025

- Identified critical limitations in existing ACH fraud detection systems where traditional single-transaction models failed to predict ACH return fraud in loyal customer segments, particularly for NSF (Non-Sufficient Funds) cases with sudden high-impact loss spikes despite strong historical performance.
- Developed novel contextual-temporal two-stage sequence modeling approach combining DistilBERT foundation model for semantic feature understanding with transformer-based temporal pattern recognition to address behavioral shift and fraud trend detection over past several transactions.
- Validated early fraud detection capability through individual driver analysis, demonstrating model correctly identified fraud pattern development before actual losses occurred while production ACH risk models showed erratic signals that failed to capture sequential fraud evolution in loyal customer segments.
- Achieved significant performance improvements on 1.52M+ production ACH transactions across 30K ACH accounts: 92% relative improvement in fraud catch rates 5% action rate, 31% reduction in false positive rates, and prevented \$1.25M additional net losses with \$3.12M gross loss reduction under \$35M Total Payment Volume compared to production ACH risk models, demonstrating superior loss prevention across all operational thresholds.

RESEARCH EXPERIENCES

Adaptive LoRA Experts Allocation and Selection for Federated Fine-Tuning

NeurIPS 2025

- Proposed FedLEASE, a novel framework addressing two critical challenges in federated LoRA fine-tuning: determining optimal number and allocation of LoRA experts across heterogeneous clients, and enabling adaptive expert selection based on client-specific data characteristics.
- Implemented data-driven clustering approach using silhouette coefficient and cosine similarity of LoRA matrices to identify optimal expert allocation, combined with innovative adaptive top-M MoE mechanism for dynamic expert selection.
- Achieved significant performance improvements on GLUE and FLAN benchmarks with RoBERTa-Large and LLaMA-2-7B models, demonstrating 5.53% average improvement over strongest baselines while maintaining communication efficiency in heterogeneous federated settings.

Federated Elastic Learning for Heterogeneous Devices

NeurIPS 2025

- Presented FedEL, a federated elastic learning framework that addresses computational heterogeneity in FL through window-based training and tensor importance adjustment mechanisms to enhance training efficiency while maintaining accuracy.
- Implemented sliding window training approach that divides DNN models into blocks and progressively trains different portions across FL rounds, combined with adaptive tensor selection using ElasticTrainer to ensure all clients complete training within similar timeframes regardless of hardware capabilities.
- Achieved up to 3.87 \times improvement in time-to-accuracy compared to FedAvg while maintaining comparable or superior final test accuracy across multiple datasets (CIFAR10, Tiny ImageNet, Google Speech Commands, Reddit) and models (VGG16, ResNet50, ALBERT), with significant reductions in memory overhead and energy consumption.

Federated LoRA Fine-Tuning with Aggregation and Initialization Refinement

ICCV 2025

- Developed LoRA-FAIR, an innovative solution to address two key challenges in federated fine-tuning of foundation models with LoRA: server-side aggregation bias and client-side initialization drift.
- Designed a residual-based correction mechanism to enhance LoRA module aggregation on the server, ensuring global updates closely approximate ideal model parameters.
- Evaluated the approach on ViT and MLP-Mixer foundation models using diverse and challenging non-IID datasets (DomainNet and NICO++), achieving up to 5.03% accuracy improvement and enhanced training stability in both feature and label non-IID settings compared to baseline methods.

Taming Cross-Domain Representation Variance in Federated Prototype Learning

NeurIPS 2024

- Designed FedPLVM, a dual-level prototype clustering algorithm with a novel α -sparsity prototype loss to address performance gaps in federated prototype learning across clients with heterogeneous data domains.
- Implemented the algorithm in Python using PyTorch, employing FINCH and cosine similarity metrics to ensure efficient training and model deployment across multiple environments.
- Achieved up to 9.88% higher accuracy on domain-shift benchmarks such as Digit-5, Office-10, and DomainNet, outperforming SOTA methods while optimizing communication efficiency and enhancing data privacy.

TECHNICAL SKILLS

General Skills: Machine Learning, Federated Learning, Computer Vision, Data Structure, LLM Fine-tuning

Languages/Frameworks: Python, Java, MATLAB, SQL, PyTorch, Keras, TensorFlow, Linux